

SVKM's NMIMS
Mukesh Patel School of Technology Management & Engineering
Program: B.Tech\MBA.Tech\MBA.Tech AI

Course: Machine Learning
Experiment No.02

PART A

(PART A : TO BE REFFERED BY STUDENTS)

A.1 Aim: To understand and implement data exploration techniques using Pandas Library.

Task 1: Perform Exploratory data analysis on Car dataset and write the inferences for each question.

- i. Read the Toyota.csv file into a DataFrame.
- ii. Explore size, shape, data types of each column in the dataset.
- iii. List down the columns of dataset
- iv. Find out 'Fuel Type' for the 4th row.
- v. Find out value for second column for the 4th row.
- vi. Select all rows for column "Fuel Type"
- vii. Select all rows for columns "KM", "HP" and "Automatic"
- viii. Display 1 to 5 rows for columns 2 to 4 (excluding row 5 and column 4)
- ix. Display the info of dataset and state your observations
- x. Identify unique values for columns "KM", "HP" and "Doors"
- xi. Create a new data frame, by replacing "?" with NAN
- xii. Replace the categorical values in the "Doors" column with its corresponding numeric value
- xiii. Convert data types of columns "Doors", "MetColor" and "Automatic" to int, and object
- xiv. Identify the total number of null values in each column of the data set
- xv. Drop rows with null values
- xvi. Identify total number of cars that runs on "Petrol", "Diesel" or "CNG"
- xvii. Identify mean of "KM" for the cars that runs on "Diesel"

Task 2:

Perform one hot and label encoding on relationship column of "adults" dataset

A.2 Prerequisite:

Python Programming, Pandas library

A.3 Outcome:

After successful completion of this experiment students will be able to:

- i. Read different types of data files(csv, excel, text file etc.)
- ii. Obtain metadata of given dataset
- iii. Understand finding of null values and replacing null values
- iv. Understand and implement class label encoding
- v. Understand and implement one hot encoding

A.4 Theory:

Exploratory Data Analysis:

Exploratory Data Analysis (EDA) is an open-ended process where we calculate statistics and make figures to find trends, anomalies, patterns, or relationships within the data. The goal of EDA is to learn what our data can tell us. It generally starts out with a high level overview, then narrows in to specific areas as we find intriguing areas of the data. The findings may be interesting in their own right, or they can be used to inform our modeling choices, such as by helping us decide which features to use.

Pandas Library:

DataFrame is a 2-dimensional labeled data structure with columns of potentially different types. You can think of it like a spreadsheet or SQL table, or a dict of Series objects. It is generally the most commonly used pandas object. Like Series, DataFrame accepts many different kinds of input:

- Dict of 1D ndarrays, lists, dicts, or Series
- 2-D numpy.ndarray
- Structured or record ndarray
- A Series
- Another DataFrame

Encoding:

One hot encoding:

One-hot encoding converts the categorical data into numeric data by splitting the column into multiple columns. The numbers are replaced by 1s and 0s, depending on which column has what value

Label encoding:

This approach is very simple and it involves converting each value in a column into a number.

PART B

(PART B : TO BE COMPLETED BY STUDENTS)

(Students must submit the soft copy as per following segments within two hours of the practical.)

Roll No. I066	Name: Srihari Thyagarajan
Class: B Tech Artificial Intelligence	Batch: B2
Date of Experiment: 22/12/2022	Date of Submission:
Grade:	

B.1 Task1

```
# ML Practical Experiment 2

[ ] # import libraries
import pandas as pd
import numpy as np
import statistics as statistics
from sklearn.preprocessing import OneHotEncoder # For Task 2
```

Task 1

```
[ ] df = pd.read_excel("/content/Toyota.csv")
```

df

	Price	Age	KM	FuelType	HP	MetColor	Automatic	CC	Doors	Weight
0	13500	23.0	46986	Diesel	90	1.0	0	2000	three	1165
1	13750	23.0	72937	Diesel	90	1.0	0	2000	3	1165
2	13950	24.0	41711	Diesel	90	NaN	0	2000	3	1165
3	14950	26.0	48000	Diesel	90	0.0	0	2000	3	1165
4	13750	30.0	38500	Diesel	90	0.0	0	2000	3	1170
...
1431	7500	NaN	20544	Petrol	86	1.0	0	1300	3	1025
1432	10845	72.0	??	Petrol	86	0.0	0	1300	3	1015
1433	8500	NaN	17016	Petrol	86	0.0	0	1300	3	1015
1434	7250	70.0	??	NaN	86	1.0	0	1300	3	1015
1435	6950	76.0	1	Petrol	110	0.0	0	1600	5	1114

1436 rows x 10 columns

```
✓ [66] # Size  
      df.size
```

14360

```
✓ [67] # Shape  
      df.shape
```

(1436, 10)

```
▶ # Data Types  
   df.dtypes
```

```
☞ Price      int64  
   Age       float64  
   KM        object  
   FuelType   object  
   HP        object  
   MetColor   object  
   Automatic  object  
   CC         int64  
   Doors      int64  
   Weight     int64  
   dtype: object
```

```
✓ [69] # Columns of a Dataset  
      for column in df.columns:  
          print(column)
```

Price
Age
KM
FuelType
HP
MetColor
Automatic
CC
Doors
Weight

```
✓ [70] # Fuel Type of the 4th row  
      df['FuelType'][3]
```

'Diesel'

```
✓ [85] # Value for second column for the 4th row  
      df.iloc[:, 2][4]
```

38500

```
38500
✓ 0s df['FuelType']
↳ 0 Diesel
   1 Diesel
   2 Diesel
   3 Diesel
   4 Diesel
   ...
 1431 Petrol
 1432 Petrol
 1433 Petrol
 1434 0
 1435 Petrol
Name: FuelType, Length: 1436, dtype: object
```

```
✓ 0s df[["FuelType", "KM", "HP"]]
↳
```

	FuelType	KM	HP
0	Diesel	46986	90
1	Diesel	72937	90
2	Diesel	41711	90
3	Diesel	48000	90
4	Diesel	38500	90
...
1431	Petrol	20544	86
1432	Petrol	NaN	86
1433	Petrol	17016	86
1434	0	NaN	86
1435	Petrol	1	110

1436 rows x 3 columns

```
✓ [130] # Value for 1-5 rows and 2-4 columns exluding the 5th row and 4th column.
      df.iloc[1: 5, 2 : 4]
```

	fnlwgt	education
1	89814	HS-grad
2	336951	Assoc-acdm
3	160323	Some-college
4	103497	Some-college

```
# Info of dataset:
df
```

	Price	Age	KM	FuelType	HP	MetColor	Automatic	CC	Doors	Weight
0	13500	23.0	46986	Diesel	90	1.0	0	2000	3	1165
1	13750	23.0	72937	Diesel	90	1.0	0	2000	3	1165
2	13950	24.0	41711	Diesel	90	0.0	0	2000	3	1165
3	14950	26.0	48000	Diesel	90	0.0	0	2000	3	1165
4	13750	30.0	38500	Diesel	90	0.0	0	2000	3	1170
...
1431	7500	0.0	20544	Petrol	86	1.0	0	1300	3	1025
1432	10845	72.0	NaN	Petrol	86	0.0	0	1300	3	1015
1433	8500	0.0	17016	Petrol	86	0.0	0	1300	3	1015
1434	7250	70.0	NaN	0	86	1.0	0	1300	3	1015
1435	6950	76.0	1	Petrol	110	0.0	0	1600	5	1114

1436 rows × 10 columns

Observations from the Dataset:

From the Dataset we observe that:

1. The columns KM -> Kilometres and Doors should have the Integer datatype. However from the dataframe we observe that some values in these columns have non-integer values.
2. The datatypes of these 2 columns have the "object" datatype.

```
df["KM"].unique()
```

```
array([46986, 72937, 41711, ..., 30964, 20544, 17016], dtype=object)
```

```
df["HP"].unique()
```

```
array([90, '????', 192, 110, 97, 71, 116, 98, 69, 86, 72, 107, 73],
      dtype=object)
```

```
df["Doors"].unique()
```

```
array(['three', 3, 5, 4, 'four', 'five', 2], dtype=object)
```

```
df[["KM", "HP", "Doors"]].nunique()
```

```
KM      1256
HP       13
Doors     7
dtype: int64
```

```
df = df.fillna(0)
df.replace('??', 'NaN', inplace = True)

# Forming a new dataframe
newdf = df.replace(to_replace = ["??", "????"], value = "NaN")
newdf

# df_1 = df.fillna(0)
# df_1.replace('??', 'NaN', inplace = True)
```

	Price	Age	KM	FuelType	HP	MetColor	Automatic	CC	Doors	Weight
0	13500	23.0	46986	Diesel	90	1.0	0	2000	three	1165
1	13750	23.0	72937	Diesel	90	1.0	0	2000	3	1165
2	13950	24.0	41711	Diesel	90	0.0	0	2000	3	1165
3	14950	26.0	48000	Diesel	90	0.0	0	2000	3	1165
4	13750	30.0	38500	Diesel	90	0.0	0	2000	3	1170
...
1431	7500	0.0	20544	Petrol	86	1.0	0	1300	3	1025
1432	10845	72.0	NaN	Petrol	86	0.0	0	1300	3	1015
1433	8500	0.0	17016	Petrol	86	0.0	0	1300	3	1015
1434	7250	70.0	NaN	0	86	1.0	0	1300	3	1015
1435	6950	76.0	1	Petrol	110	0.0	0	1600	5	1114

1436 rows × 10 columns

```
[211] # New dataframe containing ?? replaced with NaN
newdf

    Price  Age   KM  FuelType  HP  MetColor  Automatic  CC  Doors  Weight
0  13500  23.0  46986   Diesel   90      1.0         0  2000   three   1165
1  13750  23.0  72937   Diesel   90      1.0         0  2000     3    1165
2  13950  24.0  41711   Diesel   90      0.0         0  2000     3    1165
3  14950  26.0  48000   Diesel   90      0.0         0  2000     3    1165
4  13750  30.0  38500   Diesel   90      0.0         0  2000     3    1170
...
1431   7500   0.0  20544   Petrol   86      1.0         0  1300     3    1025
1432  10845  72.0   NaN   Petrol   86      0.0         0  1300     3    1015
1433   8500   0.0  17016   Petrol   86      0.0         0  1300     3    1015
1434   7250  70.0   NaN     0    86      1.0         0  1300     3    1015
1435   6950  76.0     1   Petrol  110      0.0         0  1600     5    1114

1436 rows x 10 columns

[220] # Categorical
newdf["Doors"].replace(["three", "four", "five"], [3, 4, 5], inplace = True)
```

```
[221] newdf

    Price  Age   KM  FuelType  HP  MetColor  Automatic  CC  Doors  Weight
0  13500  23.0  46986   Diesel   90      1.0         0  2000     3    1165
1  13750  23.0  72937   Diesel   90      1.0         0  2000     3    1165
2  13950  24.0  41711   Diesel   90      0.0         0  2000     3    1165
3  14950  26.0  48000   Diesel   90      0.0         0  2000     3    1165
4  13750  30.0  38500   Diesel   90      0.0         0  2000     3    1170
...
1431   7500   0.0  20544   Petrol   86      1.0         0  1300     3    1025
1432  10845  72.0   NaN   Petrol   86      0.0         0  1300     3    1015
1433   8500   0.0  17016   Petrol   86      0.0         0  1300     3    1015
1434   7250  70.0   NaN     0    86      1.0         0  1300     3    1015
1435   6950  76.0     1   Petrol  110      0.0         0  1600     5    1114

1436 rows x 10 columns

newdf["Doors"] = newdf["Doors"].astype(int)
newdf["MetColor"] = newdf["MetColor"].astype(object)
newdf["Automatic"] = newdf["Automatic"].astype(object)
```

```
[223] df_1 = pd.read_excel("/content/Toyota.csv")
df_1.isnull().sum()

Price      0
Age       100
KM         0
FuelType   100
HP         0
MetColor   150
Automatic   0
CC         0
Doors       0
Weight     0
dtype: int64
```

```
newdf.dropna()

Price Age KM FuelType HP MetColor Automatic CC Doors Weight
0 13500 23.0 46986 Diesel 90 1.0 0 2000 3 1165
1 13750 23.0 72937 Diesel 90 1.0 0 2000 3 1165
2 13950 24.0 41711 Diesel 90 0.0 0 2000 3 1165
3 14950 26.0 48000 Diesel 90 0.0 0 2000 3 1165
4 13750 30.0 38500 Diesel 90 0.0 0 2000 3 1170
... ... ... ... ... ... ... ... ... ...
1431 7500 0.0 20544 Petrol 86 1.0 0 1300 3 1025
1432 10845 72.0 NaN Petrol 86 0.0 0 1300 3 1015
1433 8500 0.0 17016 Petrol 86 0.0 0 1300 3 1015
1434 7250 70.0 NaN 0 86 1.0 0 1300 3 1015
1435 6950 76.0 1 Petrol 110 0.0 0 1600 5 1114

1436 rows x 10 columns

[225] newdf["FuelType"].value_counts()

FuelType
Petrol 1177
Diesel 144
0 100
CNG 15
Name: FuelType, dtype: int64
```

```
[218] newdf

[226] l = []

for i in range(len(newdf['FuelType'])):
    if newdf['FuelType'][i] == 'Diesel':
        if newdf['KM'][i] != 'NaN':
            l.append(int(newdf['KM'][i]))

np.mean(l)

114927.87857142858
```


B.2 Task 2

Task 2

```
[120] df = pd.read_excel("/content/adult.csv")
      df_1 = pd.read_excel("/content/adult.csv")
```

```
[123] # Checking for the labels in the categorical parameters
      df_1["relationship"].unique()

      array(['Own-child', 'Husband', 'Not-in-family', 'Unmarried', 'Wife',
            'Other-relative'], dtype=object)
```

```
[124] # Checking for the label counts in the categorical parameters
      df_1["relationship"].value_counts()
```

```
Husband      19716
Not-in-family 12583
Own-child     7581
Unmarried     5125
Wife          2331
Other-relative 1506
Name: relationship, dtype: int64
```

Method 1:

One Hot Encoding using Sci-kit learn Library:

```
[110] # Creating an instance of the one-hot-encoder
      encoder = OneHotEncoder(handle_unknown='ignore')

      # Perform one-hot encoding on 'relationship' column
      encoder_df = pd.DataFrame(encoder.fit_transform(df[['relationship']]).toarray())
```

```
[111] # Merging one-hot encoded columns back with original DataFrame df.
      final_df = df.join(encoder_df)
```

	age	workclass	fnlwgt	education	educational-num	marital-status	occupation	relationship	race	gender	...	capital-loss	hours-per-week	native-country	income	0	1	2	3	4
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-child	Black	Male	...	0	40	United-States	<=50K	0.0	0.0	0.0	1.0	0.0
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	White	Male	...	0	50	United-States	<=50K	1.0	0.0	0.0	0.0	0.0
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband	White	Male	...	0	40	United-States	>50K	1.0	0.0	0.0	0.0	0.0
3	44	Private	160323	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband	Black	Male	...	0	40	United-States	>50K	1.0	0.0	0.0	0.0	0.0
4	18	?	103497	Some-college	10	Never-married	?	Own-child	White	Female	...	0	30	United-States	<=50K	0.0	0.0	0.0	1.0	0.0
...
48837	27	Private	257302	Assoc-acdm	12	Married-civ-spouse	Tech-support	Wife	White	Female	...	0	38	United-States	<=50K	0.0	0.0	0.0	0.0	1.0
48838	40	Private	154374	HS-grad	9	Married-civ-spouse	Machine-op-inspct	Husband	White	Male	...	0	40	United-States	>50K	1.0	0.0	0.0	0.0	0.0
48839	58	Private	151910	HS-grad	9	Widowed	Adm-clerical	Unmarried	White	Female	...	0	40	United-States	<=50K	0.0	0.0	0.0	0.0	1.0
48840	22	Private	201490	HS-grad	9	Never-married	Adm-clerical	Own-child	White	Male	...	0	20	United-States	<=50K	0.0	0.0	0.0	1.0	0.0
48841	52	Self-emp-inc	287927	HS-grad	9	Married-civ-spouse	Exec-managerial	Wife	White	Female	...	0	40	United-States	>50K	0.0	0.0	0.0	0.0	1.0

48842 rows x 21 columns

```
[113] # Dropping the original relationship column from the dataframe as we will be referring only to the Numerical values
# which are to be generated

final_df.drop('relationship', axis=1, inplace=True)
```

final_df

	age	workclass	fnlwt	education	educational-num	marital-status	occupation	race	gender	capital-gain	capital-loss	hours-per-week	native-country	income	0	1	2	3	4	5
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Black	Male	0	0	40	United-States	<=50K	0.0	0.0	0.0	1.0	0.0	0.0
1	38	Private	89814	HS-grad	9	Married-div-spouse	Farming-fishing	White	Male	0	0	50	United-States	<=50K	1.0	0.0	0.0	0.0	0.0	0.0
2	28	Local-gov	336951	Assoc-acdm	12	Married-div-spouse	Protective-serv	White	Male	0	0	40	United-States	>50K	1.0	0.0	0.0	0.0	0.0	0.0
3	44	Private	160323	Some-college	10	Married-div-spouse	Machine-op-inspct	Black	Male	7688	0	40	United-States	>50K	1.0	0.0	0.0	0.0	0.0	0.0
4	18	?	103497	Some-college	10	Never-married	?	White	Female	0	0	30	United-States	<=50K	0.0	0.0	0.0	1.0	0.0	0.0
...
48837	27	Private	257302	Assoc-acdm	12	Married-div-spouse	Tech-support	White	Female	0	0	38	United-States	<=50K	0.0	0.0	0.0	0.0	0.0	1.0
48838	40	Private	154374	HS-grad	9	Married-div-spouse	Machine-op-inspct	White	Male	0	0	40	United-States	>50K	1.0	0.0	0.0	0.0	0.0	0.0
48839	58	Private	151910	HS-grad	9	Widowed	Adm-clerical	White	Female	0	0	40	United-States	<=50K	0.0	0.0	0.0	0.0	1.0	0.0
48840	22	Private	201490	HS-grad	9	Never-married	Adm-clerical	White	Male	0	0	20	United-States	<=50K	0.0	0.0	0.0	1.0	0.0	0.0
48841	52	Self-emp-inc	287927	HS-grad	9	Married-div-spouse	Exec-managerial	White	Female	15024	0	40	United-States	>50K	0.0	0.0	0.0	0.0	0.0	1.0

```
[126] final_df.columns = ['age', 'workclass', 'fnlwt', 'education', 'educational-num', 'marital-status', 'occupation', 'race', 'gender', 'capital-gain', 'capital-loss', 'hours-per-week', 'native-country', 'income', 'Own-child', 'Husband', 'Not-in-family', 'Unmarried', 'Wife', 'Other-relative']
```

	age	workclass	fnlwt	education	educational-num	marital-status	occupation	race	gender	capital-gain	capital-loss	hours-per-week	native-country	income	Own-child	Husband	Not-in-family	Unmarried	Wife	Other-relative
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Black	Male	0	0	40	United-States	<=50K	0.0	0.0	0.0	1.0	0.0	0.0
1	38	Private	89814	HS-grad	9	Married-div-spouse	Farming-fishing	White	Male	0	0	50	United-States	<=50K	1.0	0.0	0.0	0.0	0.0	0.0
2	28	Local-gov	336951	Assoc-acdm	12	Married-div-spouse	Protective-serv	White	Male	0	0	40	United-States	>50K	1.0	0.0	0.0	0.0	0.0	0.0
3	44	Private	160323	Some-college	10	Married-div-spouse	Machine-op-inspct	Black	Male	7688	0	40	United-States	>50K	1.0	0.0	0.0	0.0	0.0	0.0
4	18	?	103497	Some-college	10	Never-married	?	White	Female	0	0	30	United-States	<=50K	0.0	0.0	0.0	1.0	0.0	0.0
...
48837	27	Private	257302	Assoc-acdm	12	Married-div-spouse	Tech-support	White	Female	0	0	38	United-States	<=50K	0.0	0.0	0.0	0.0	0.0	1.0
48838	40	Private	154374	HS-grad	9	Married-div-spouse	Machine-op-inspct	White	Male	0	0	40	United-States	>50K	1.0	0.0	0.0	0.0	0.0	0.0
48839	58	Private	151910	HS-grad	9	Widowed	Adm-clerical	White	Female	0	0	40	United-States	<=50K	0.0	0.0	0.0	0.0	1.0	0.0

Method 2

One-Hot encoding the categorical parameters using get_dummies()

```
[128] one_hot_encoded_data = pd.get_dummies(df_1, columns = ['relationship'])
```

one_hot_encoded_data																			
	age	workclass	fnlwt	education	educational-num	marital-status	occupation	race	gender	capital-gain	capital-loss	hours-per-week	native-country	income	relationship_Husband	relationship_Wife	relationship_in-family	relationship	relationship
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Black	Male	0	0	40	United-States	<=50K	0	0	0	0	0
1	38	Private	89614	HS-grad	9	Married-civ-spouse	Farming-fishing	White	Male	0	0	50	United-States	<=50K	1	0	0	0	0
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	Protective-serv	White	Male	0	0	40	United-States	>50K	1	0	0	0	0
3	44	Private	160323	Some-college	10	Married-civ-spouse	Machine-op-inspct	Black	Male	7668	0	40	United-States	>50K	1	0	0	0	0
4	18	?	103497	Some-college	10	Never-married	?	White	Female	0	0	30	United-States	<=50K	0	0	0	0	0
...
48837	27	Private	257302	Assoc-acdm	12	Married-civ-spouse	Tech-support	White	Female	0	0	38	United-States	<=50K	0	0	0	0	0
48838	40	Private	154374	HS-grad	9	Married-civ-spouse	Machine-op-inspct	White	Male	0	0	40	United-States	>50K	1	0	0	0	0
48839	58	Private	151910	HS-grad	9	Widowed	Adm-clerical	White	Female	0	0	40	United-States	<=50K	0	0	0	0	0
48840	22	Private	201490	HS-grad	9	Never-married	Adm-clerical	White	Male	0	0	20	United-States	<=50K	0	0	0	0	0
48841	52	Self-emp-inc	267927	HS-grad	9	Married-civ-spouse	Exec-managerial	White	Female	15024	0	40	United-States	>50K	0	0	0	0	0

48842 rows x 20 columns

Warning: total number of rows (48842) exceeds max_rows (20000). Limiting to first (20000) rows.

B.4 Conclusion:

From the above experiment, I learnt the following:

- Read different types of data files (csv, excel, text file etc.).
- Obtain metadata of given dataset.
- Understand finding of null values and replacing null values.
- Understand and implement class label encoding.
- Understand and implement one hot encoding.